



4주차 머신러닝 발표

김찬원

Convolutional Neural Network

- Most widely used for image classification.
- Generally, it consists of convolution layer, pooling layer and fully-connected layer.
- Convolution, Pooling layer – feature extraction
- Fully-connected layer – classification





Output size:
 $(N - F) / \text{stride} + 1$

e.g. $N = 7, F = 3$:

stride 1 $\Rightarrow (7 - 3) / 1 + 1 = 5$

stride 2 $\Rightarrow (7 - 3) / 2 + 1 = 3$

stride 3 $\Rightarrow (7 - 3) / 3 + 1 = 2.33 : \backslash$

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!


in general, common to see CONV layers with stride 1, filters of size $F \times F$, and zero-padding with $(F-1)/2$. (will preserve size spatially)

e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3

Swiping the entire image




$32 \times 32 \times 3$ image

activation maps
(?, ?, ?)


2D Convolution Layer

– Multi Channel, Many Filters




Pooling

- Max Pooling or Average Pooling



Fully Connected(Dense) Layer



Case Study: AlexNet

[Krizhevsky et al. 2012]

Full (simplified) AlexNet architecture:

[227x227x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1


[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)



Details/Retrospectives:

- first use of ReLU
- used Norm layers (not common anymore)
- heavy data augmentation
- dropout 0.5
- batch size 128
- SGD Momentum 0.9
- Learning rate $1e-2$, reduced by 10 manually when val accuracy plateaus
- L2 weight decay $5e-4$
- 7 CNN ensemble: 18.2% -> 15.4%

tf.keras.layers.Conv2D


- **filters** : Integer, the dimensionality of the output space (i.e. the number of output filters in the convolution).
- **kernel_size** : An integer or tuple/list of 2 integers, specifying the height and width of the 2D convolution window. Can be a single integer to specify the same value for all spatial dimensions.
- **strides** : An integer or tuple/list of 2 integers, specifying the strides of the convolution along the height and width. Can be a single integer to specify the same value for all spatial dimensions. Specifying any stride value != 1 is incompatible with specifying any **dilation_rate** value != 1.
- **padding** : one of "valid" or "same" (case-insensitive).
- **data_format** : A string, one of **channels_last** (default) or **channels_first**. The ordering of the dimensions in the inputs. **channels_last** corresponds to inputs with shape (batch, height, width, channels) while **channels_first** corresponds to inputs with shape (batch, channels, height, width). It defaults to the **image_data_format** value found in your Keras config file at ~/.keras/keras.json. If you never set it, then it will be "channels_last".

tf.keras.layers.Conv2D

- **activation** : Activation function to use. If you don't specify anything, no activation is applied (ie. "linear" activation: $a(x) = x$).
- **use_bias** : Boolean, whether the layer uses a bias vector.
- **kernel_initializer** : Initializer for the **kernel** weights matrix.
- **bias_initializer** : Initializer for the bias vector.
- **kernel_regularizer** : Regularizer function applied to the **kernel** weights matrix.
- **bias_regularizer** : Regularizer function applied to the bias vector.

kernel dimension : {height, width, in_channel, out_channel}

Ex) {5, 5, 3, 2}



NN Implementation Flow in TensorFlow

1. Set hyper parameters – learning rate, training epochs, batch size, etc.
2. Make a data pipelining – use `tf.data`
3. Build a neural network model – use `tf.keras` sequential APIs
4. Define a loss function – cross entropy
5. Calculate a gradient – use `tf.GradientTape`
6. Select an optimizer – Adam optimizer
7. Define a metric for model's performance – accuracy
8. (optional) Make a checkpoint for saving
9. Train and Validate a neural network model