

# ML/DL Study

```
lookup.KeyValue  
f.constant(['em  
=tf.constant([0  
ce = tf.lookup.StaticV  
init,  
num_oov_buckets=5)  
  
lookup.StaticVocabular  
initializer,  
num_oov_buckets,  
lookup_key_dtype=None  
name=None,  
experimental_is_open
```

# Code Review

## Code Review

# Exploring Data

```
df.head()
```

Python

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	2008	WD	Normal	208500
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	2007	WD	Normal	181500
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	2008	WD	Normal	223500
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	2006	WD	Abnorml	140000
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	2008	WD	Normal	250000

5 rows × 81 columns

```
df = df.drop('Id', axis=1)
df.head(3)
```

Python

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighborhood	Condition1	Condition2	BldgType	HouseStyle	OverallQual	OverallCond	YearBuilt
0	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	Norm	1Fam	2Story	7	5	2003
1	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	AllPub	FR2	Gtl	Veenker	Feedr	Norm	1Fam	1Story	6	8	1976
2	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	AllPub	Inside	Gtl	CollgCr	Norm	Norm	1Fam	2Story	7	5	2001

```
df = df.dropna(axis=1, how='any')
```

MSSubClass	0	MSSubClass	0
MSZoning	0	MSZoning	0
LotFrontage	259	LotArea	0
LotArea	0	Street	0
Street	0	LotShape	0
...		...	
MoSold	0	MoSold	0
YrSold	0	YrSold	0
SaleType	0	SaleType	0
SaleCondition	0	SaleCondition	0
SalePrice	0	SalePrice	0
Length: 80, dtype: int64		Length: 61, dtype: int64	

```
def min_max_scaling(data):
    min_val = data.min()
    max_val = data.max()
    scaled_data = (data - min_val) / (max_val - min_val)
    return scaled_data
```

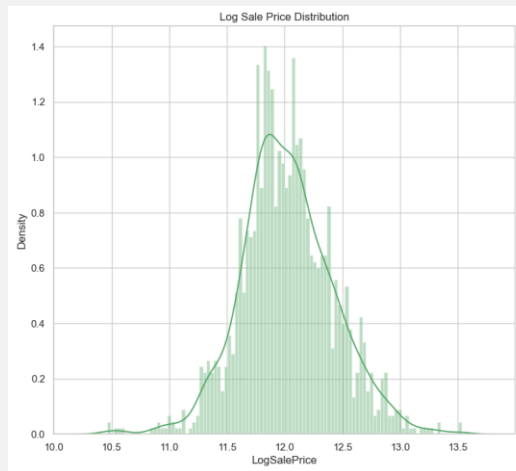

```
df['LogSalePrice'] = min_max_scaling(df['LogSalePrice'])
print(df['LogSalePrice'].describe())
```

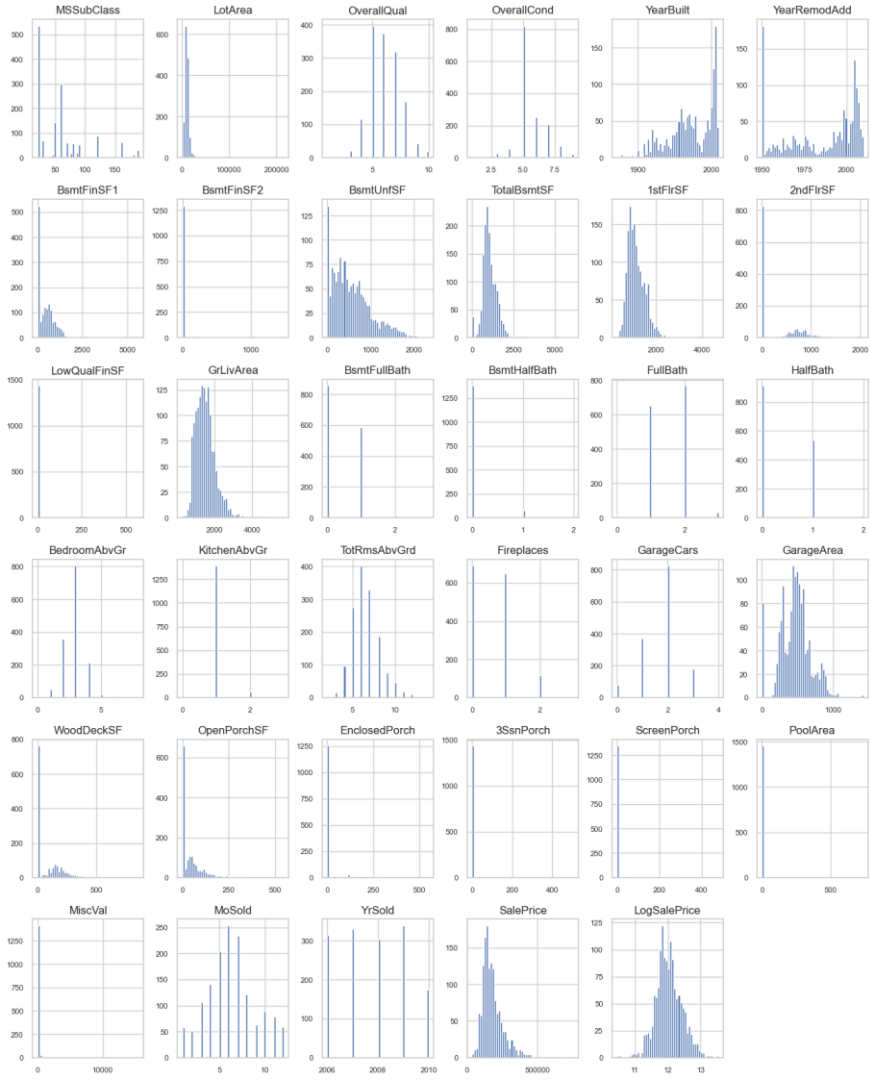
✓ 0.0s

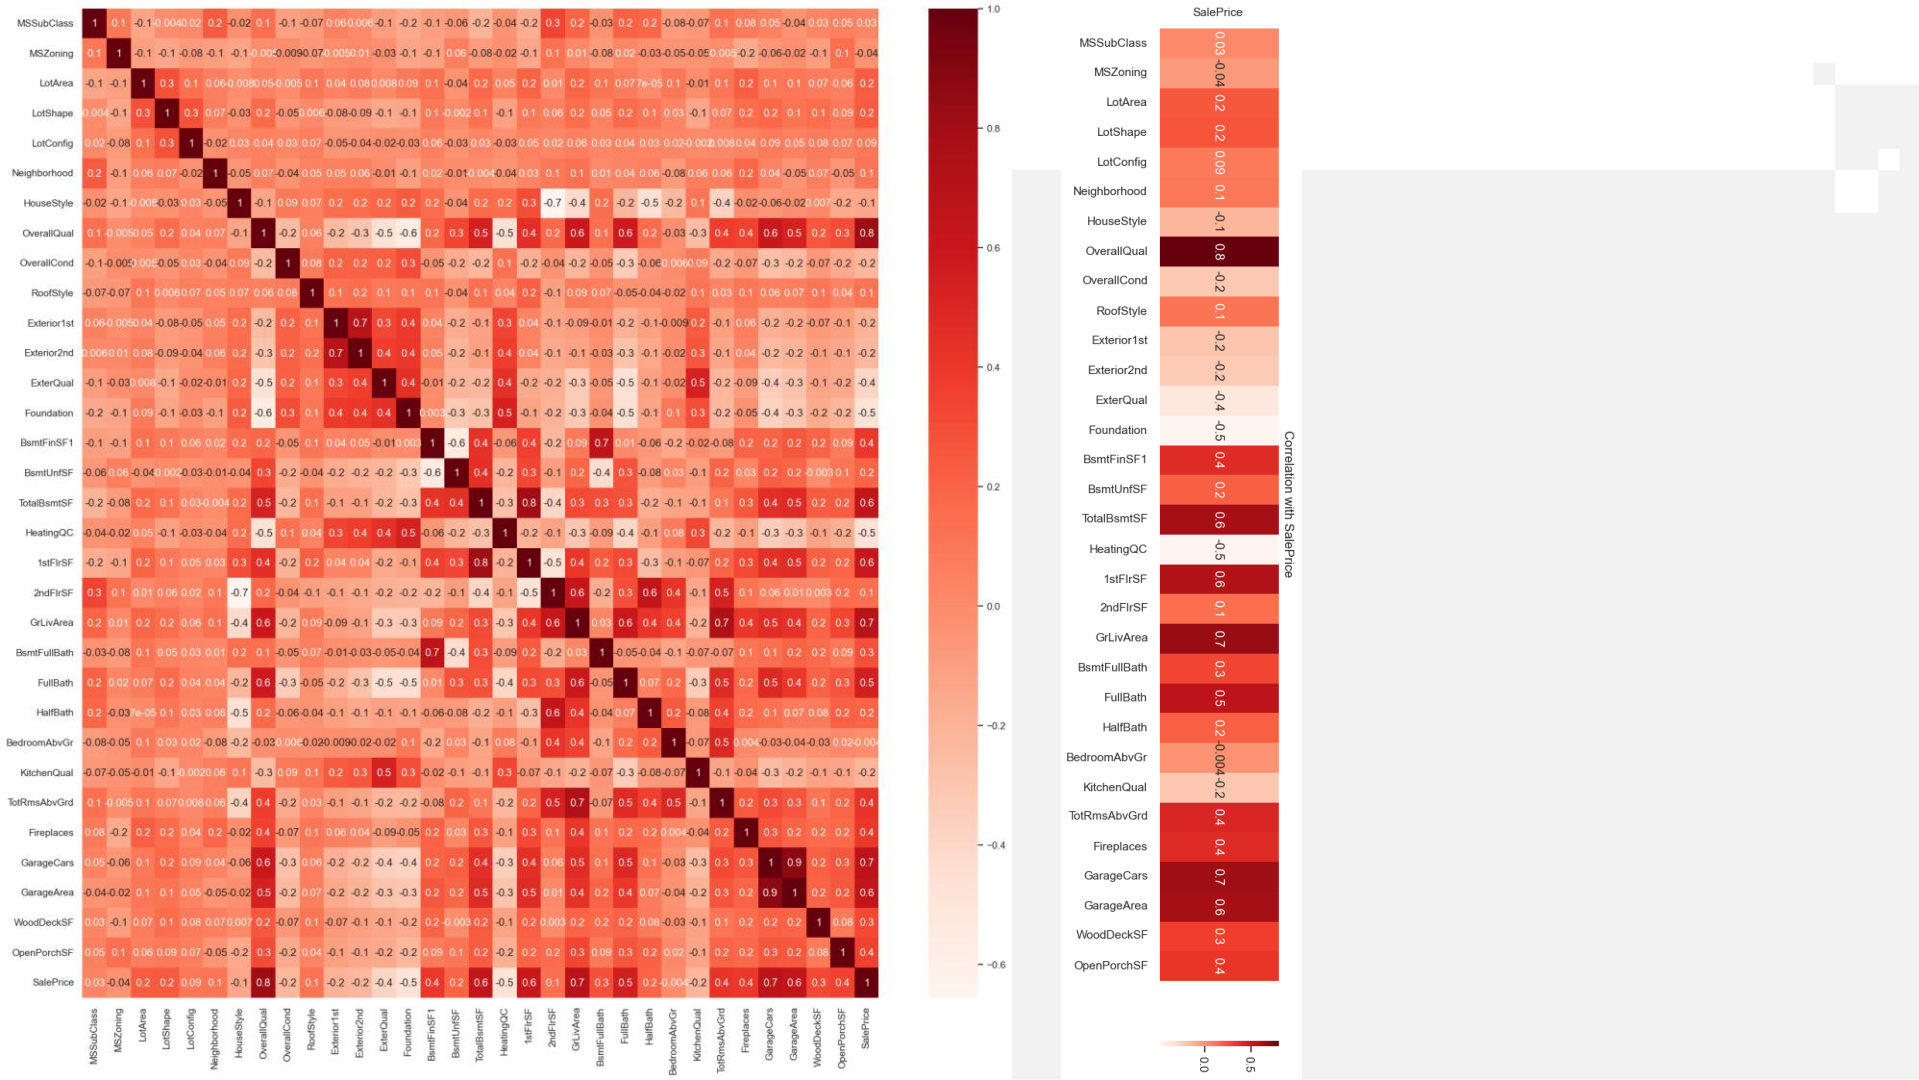
Python

```
count    1460.000000
mean      0.508683
std       0.129936
min       0.000000
25%      0.427702
50%      0.501349
75%      0.589900
max       1.000000
Name: LogSalePrice, dtype: float64
```

```
df['LogSalePrice'] = np.log(df['SalePrice'])
```







```
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_train)
mse = mean_squared_error(y_train, y_pred)
r2 = r2_score(y_train, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```

✓ 0.0s

Python

Mean Squared Error: 621206121.4955989

R-squared: 0.8460515711954664

# Random Forest

```
from sklearn.ensemble import RandomForestRegressor

model_1 = RandomForestRegressor(n_estimators=100, random_state=0)

model_1.fit(X_train, y_train)

y_pred = model_1.predict(X_train)
mse = mean_squared_error(y_train, y_pred)
r2 = r2_score(y_train, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```

✓ 3.5s

Python

Mean Squared Error: 89644858.08237426

R-squared: 0.9777840485232812



# Gradient Boosting

```
from sklearn.ensemble import GradientBoostingRegressor

model_2 = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, random_state=0)

model_2.fit(X_train, y_train)

y_pred = model_2.predict(X_train)
mse = mean_squared_error(y_train, y_pred)
r2 = r2_score(y_train, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```

✓ 0.8s

Python

Mean Squared Error: 250821838.44492674

R-squared: 0.9378408766393095

# XGB

```
import xgboost as xgb

model_4 = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1, random_state=0)

model_4.fit(X_train, y_train)

y_pred = model_4.predict(X_train)
mse = mean_squared_error(y_train, y_pred)
r2 = r2_score(y_train, y_pred)


print(f"Mean Squared Error: {mse}")
print(f"R-squared: {r2}")
```


✓ 0.2s

Python

Mean Squared Error: 45084063.98391847

R-squared: 0.9888271854151449

#	Team	Members	Score	Entries	Last	Join
3781	goldbabyerim		0.18915	1	35s	



Your First Entry!  
Welcome to the leaderboard!